



The Golden Thread

How a 40-year-old design principle is solving
the biggest challenge in modern AI.



In the beginning, computers were for experts.

The late 1970s. Computing meant command lines, memorized commands, and a high cognitive load. Systems were powerful, but not intuitive. Users bore the burden of remembering commands and managing the machine's complexity.




```
root@mainframe:~# ls -laR /var/log/
total 1432
drwxr-xr-x 2 root root 4096 Dec 12 10:01 ./
drwxr-xr-x 3 root root 4096 Dec 12 09:55 ../
-rw-r--r-- 1 root root 784 Dec 12 09:55 auth.log
-rw-r--r-- 1 root root 12381 Dec 12 10:01 debug.log
-rw-r--r-- 1 root root 9841 Dec 11 17:32 error.log
-rw-r--r-- 1 root root 2345 Dec 11 17:32 kernel.log
-rw-r--r-- 1 root root 34981 Dec 10 14:22 secure.log
-rw-r--r-- 1 root root 1245 Dec 10 14:22 syslog.log
-rw-r--r-- 1 root root 982 Dec 09 11:15 Xorg.0.log
root@mainframe:~# chmod 755 /opt/system/bin/config_tool
root@mainframe:~# grep -r 'ERROR' /var/log/ > /tmp/errors.txt
root@mainframe:~# ps -ef | grep systemd
root@mainframe:~# _
```

The dominant interaction model was 'remembering and typing.'

Then, a team at Xerox PARC asked a revolutionary question:
How can we make computers
intuitive for everyone?

Their answer was the Xerox Star, a system built not around programs, but around the user's data and mental model. It introduced foundational concepts we still use today: the Desktop Metaphor, Direct Manipulation, and Generic Commands.


The Star's philosophy was built on
“emphasizing recognition over recall,
seeing and pointing over remembering
and typing.”



Among the Star's innovations was a core principle for managing complexity: Progressive Disclosure.

The designers recognized that overwhelming users with information was as bad as making them memorize commands. Their solution was elegant and intentional.

“Progressive disclosure dictates that detail be hidden from users until they ask or need to see it.”




Detail is hidden until the user asks to see it.

The system provided default settings but also hid those settings until a user indicated they wanted to change them. This avoided overwhelming users with irrelevant choices and information.

Decades later, psychology proved the Star's designers were right. More transparency isn't always better.

We intuitively believe more information helps us make better decisions. But modern research reveals a paradox: sometimes, more detail is just noise. It can be distracting, undermine confidence, and lead us to form incorrect mental models of how a system works.




Users found that detailed, incremental feedback was "distracting and potentially undermines simple heuristics users form of system operation."

An experiment tested two interfaces: one with high transparency, one with less. Users' initial preferences were completely wrong.

- In a study, users interacted with an emotion-detecting algorithm ("E-meter").
- Word-level feedback (High Transparency): Showed how every single word influenced the result.
 - Document-level feedback (Low Transparency): Showed only the final, overall result.


High Transparency



Low Transparency



Preference Reversal



The Finding: Before using the systems, users predicted the high-transparency version would be more accurate. After using both, they were evenly split, with many preferring the simpler version.

Why?: The less transparent system hid minor errors. This allowed users to form a more positive and effective working heuristic of the system, leading some to 'overestimate system capabilities.'

Today, we face a new complexity crisis: the AI context window.




“Context quality is
the bottleneck.”

AI coding agents need to understand vast codebases, often comprising “thousands of files and millions of lines of code.” But the model’s attention—its context window—is a finite and expensive resource.

The Problem

You cannot simply feed an entire codebase into a prompt. The quality of the code an agent generates is directly determined by the relevance of the code it can “see.”



The demand for less noise is universal. Developers are literally asking for it.

On a popular open-source AI project, a user filed an enhancement request. The problem? The AI tool's output was too verbose, filled with technical details that obscured the final, important answer.

Current Behavior (Verbose)

```
Search(pattern: "optimize database query")
Read(file: "src/db/connector.py")
Analyze(code: "...")
Search(pattern: "index creation")
Read(file: "docs/performance.md")
```

Final Response: The database query can be optimized by adding an index on the 'user_id' column and refactoring the join operation.

Desired Behavior (with --quiet)

Final Response: The database query can be optimized by adding an index on the 'user_id' column and refactoring the join operation.


“Tool call output adds 6+ lines of noise before the actual response... Users must scroll past technical details to see the summary.” This is a user-driven demand for **progressive disclosure.**

The solution is a direct reincarnation of a 40-year-old principle.

An industry-wide pattern has emerged to manage AI context: **Agent Skills**. A Skill is a packaged set of instructions and tools that an AI agent can load on demand.

How it is Progressive Disclosure

1. **Initial State:** The agent only sees the `name` and `description` of all available Skills. This is a very small, low-token overview.
2. **On-Demand Loading:** When a task matches a Skill's description, the agent then chooses to load the full `SKILL.md` file, which contains detailed instructions and references to other files.



The Result: Detail is hidden from the agent until it asks for or needs to see it.

This isn't just one company's idea. It's becoming a de facto standard.

The `SKILL.md` format for packaging agent knowledge has been adopted across a range of influential coding tools.



"What surprised me is this isn't Claude-only anymore. The format got adopted by: Cursor, VS Code / GitHub Copilot, OpenAI's Codex CLI... So you can write a skill once and use it across tools. It's like how .editorconfig standardized formatting rules across editors, but for agent workflows."

The challenges change, but the principles of good design endure.



The future of AI is not about providing all the information. It's about designing for its timely and intentional disclosure. The lessons from the past don't just guide us; they provide the blueprint.